

Multiclass Classification: **One-vs-All** and **All-Pairs**

Adam Ibrahim, Nicolas Kim, Erin Okey

https://github.com/nickkim1/data_2060_final_project.git

One-vs-All

Binary learner
for each class

All-Pairs

Binary learner
for each unique pair
of classes

Implemented by
combining hypotheses
from **binary methods**

Cross-entropy loss

Optimized with
stochastic gradient
descent

Representation

Data

d numeric features

$$\mathcal{X} = \mathbb{R}^d$$

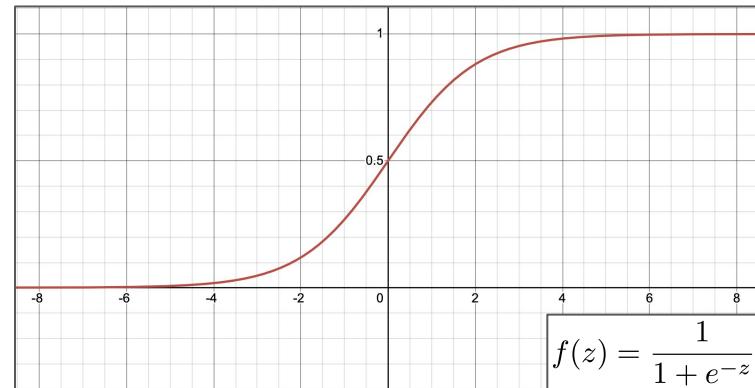
labeled 1 or 0

$$\mathcal{Y} = \{ 1, 0 \}$$

Hypothesis class

$$h_{\mathbf{w}}(\mathbf{x}) = \frac{1}{1 + e^{-\langle \mathbf{w}, \mathbf{x} \rangle}}$$

probability of class 1



Loss

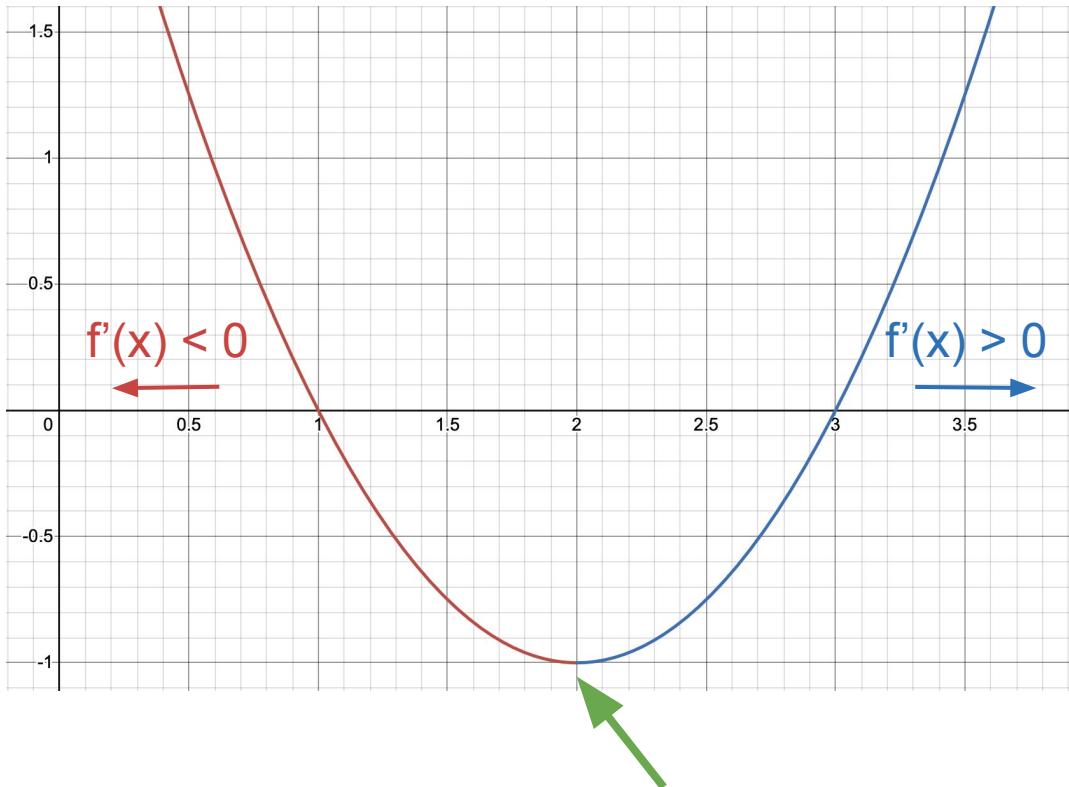
$$L_s(h_w) = -\frac{1}{m} \sum_{i=1}^m (y_i \log h_w(x_i) + (1 - y_i) \log(1 - h_w(x_i)))$$

$$\mathcal{Y} = \{ \textcolor{green}{1}, \textcolor{red}{0} \}$$

$$\begin{aligned} \log(z) &\rightarrow -\infty & \text{as } z \rightarrow 0 \\ \log(z) &\rightarrow 0 & \text{as } z \rightarrow 1 \end{aligned}$$

Optimizer - Gradient Descent

$$w \leftarrow w - \alpha \frac{\partial L(h_w)}{\partial w}$$



Optimizer - Stochastic Gradient Descent

Idea:

$$w \leftarrow w - \alpha \frac{\partial L(h_w)}{\partial w}$$

Time consuming to compute this using all data points. Instead, only use a random subset

Inputs: training examples S , step size α , batch size $b < |S|$

Initialize $w \in \mathbb{R}^{k \times d}$ randomly

converged \leftarrow False

while !converged:

 Shuffle S

for $i = 0 \dots \lceil |S| / b \rceil - 1$:

$S' \leftarrow S[i \cdot b : (i + 1) \cdot b]$

$w \leftarrow w - \alpha \nabla L_{S'}(h_w)$

 converged \leftarrow check_convergence(S, w)

return

Algorithm (Pseudocode)

One-versus-All

input:

training set $S = (\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$

algorithm for binary classification A

foreach $i \in \mathcal{Y}$

let $S_i = (\mathbf{x}_1, (-1)^{\mathbb{1}_{[y_1 \neq i]}}), \dots, (\mathbf{x}_m, (-1)^{\mathbb{1}_{[y_m \neq i]}})$

let $h_i = A(S_i)$

output:

the multiclass hypothesis defined by $h(\mathbf{x}) \in \operatorname{argmax}_{i \in \mathcal{Y}} h_i(\mathbf{x})$

All-Pairs

input:

training set $S = (\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$

algorithm for binary classification A

foreach $i, j \in \mathcal{Y}$ s.t. $i < j$

initialize $S_{i,j}$ to be the empty sequence

for $t = 1, \dots, m$

If $y_t = i$ add $(\mathbf{x}_t, 1)$ to $S_{i,j}$

If $y_t = j$ add $(\mathbf{x}_t, -1)$ to $S_{i,j}$

let $h_{i,j} = A(S_{i,j})$

output:

the multiclass hypothesis defined by

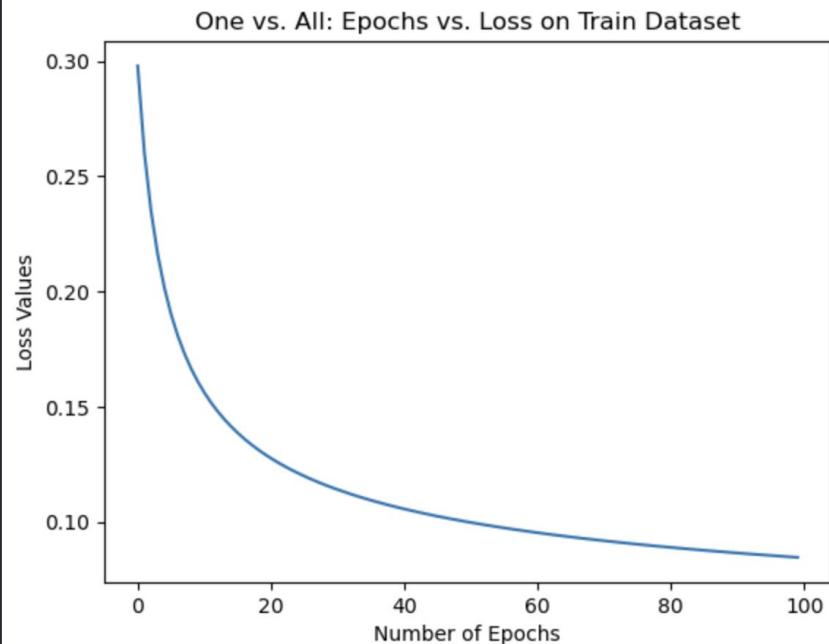
$$h(\mathbf{x}) \in \operatorname{argmax}_{i \in \mathcal{Y}} \left(\sum_{j \in \mathcal{Y}} \operatorname{sign}(j - i) h_{i,j}(\mathbf{x}) \right)$$

Hyperparameters (for both one vs. rest and all pairs)

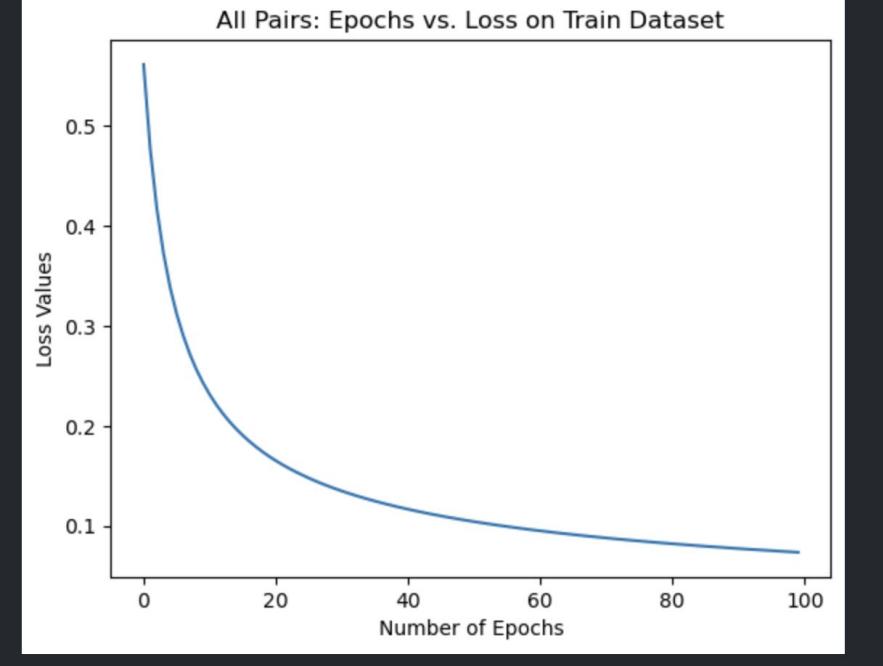
- **Dataset size:** 2000 training examples
 - Digit labels encoded as integers: 0-9 and pixel features: normalized between 0-1 (divide all values by 255)
- **Batch size:** 32
- **Lambda:** 0.001
- **Num epochs:** 100

Accuracies and Loss Curves (One-vs-All - L, All-Pairs - R)

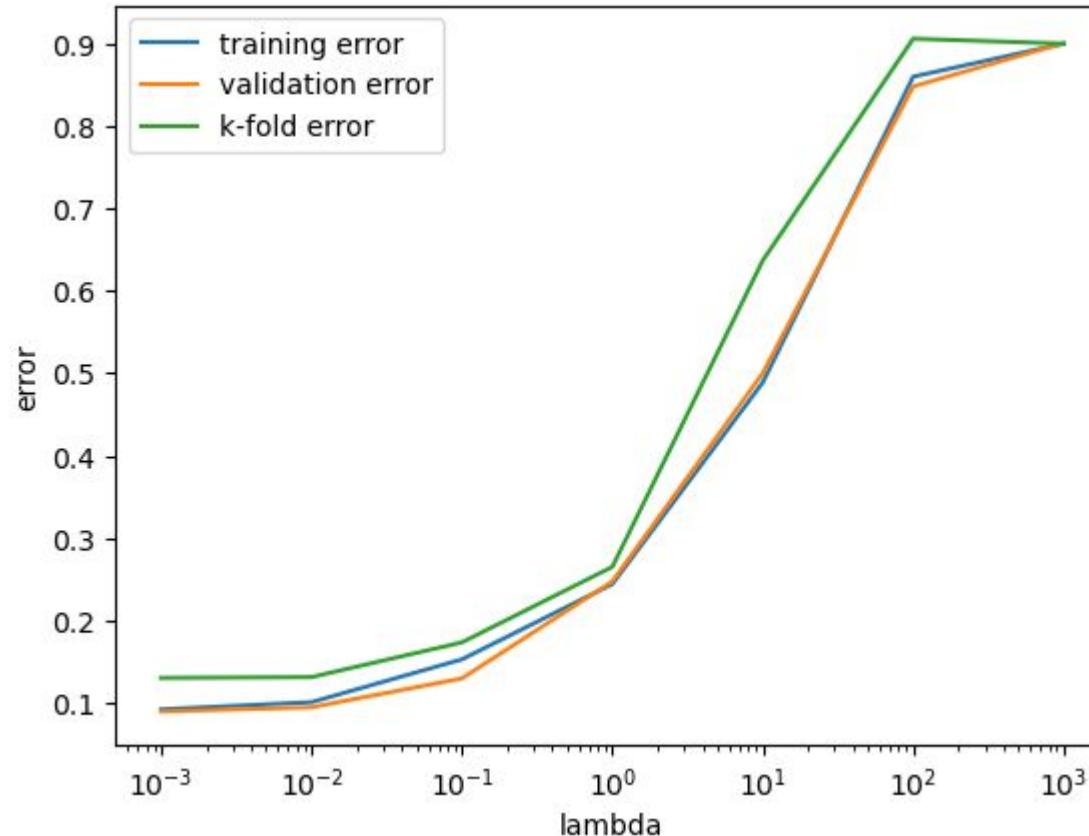
One vs. All: Train Accuracy: 0.906875
One vs. All: Validation Accuracy: 0.91



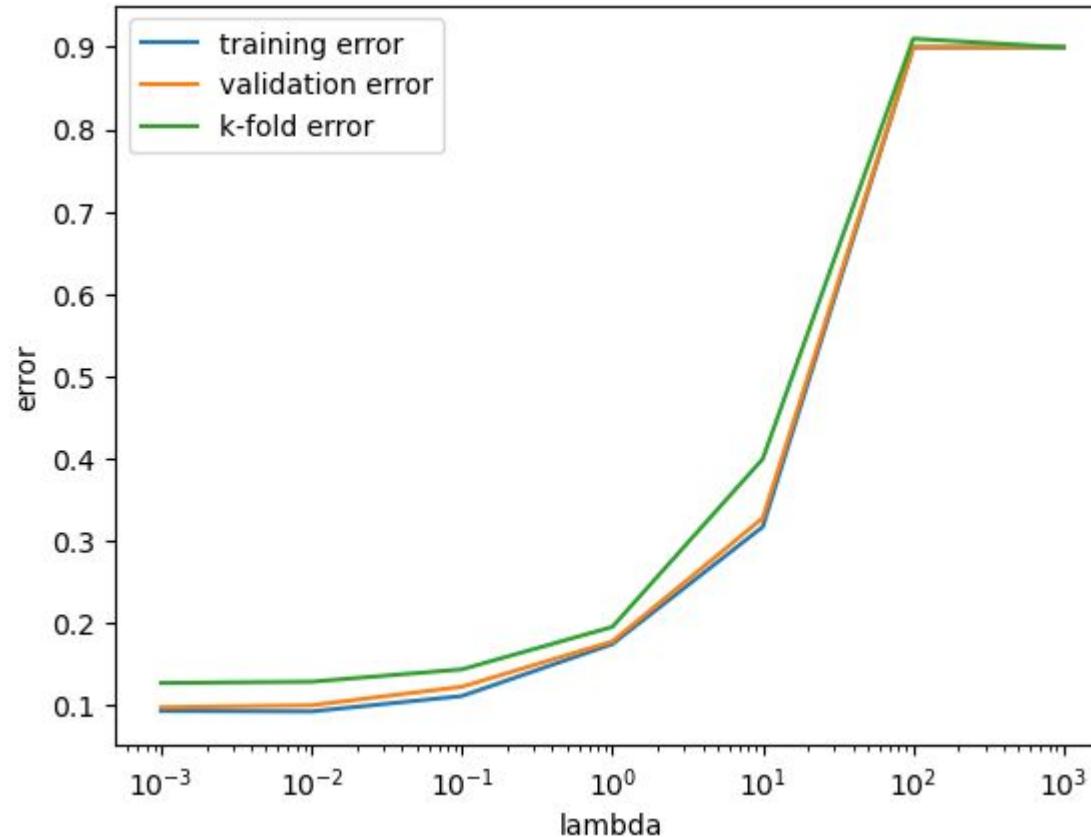
All Pairs: Train Accuracy: 0.90875
All Pairs: Validation Accuracy: 0.905



Error vs Regularization Strength (One-vs-All)



Error vs Regularization Strength (All-Pairs)



Sklearn Comparison: One-vs-All and All-Pairs Results

```
== One-vs-Rest (one vs. all) comparison ==
Our OVR logistic regression accuracy:      0.9075
Sklearn OVR logistic regression accuracy: 0.875
Predictions identical on this dataset? False
```

```
== One-vs-One (all-pairs) comparison ==
Our OVO logistic regression accuracy:      0.9025
Sklearn OVO logistic regression accuracy: 0.9075
Predictions identical on this dataset? False
```

- **Explanation of Differences:** Our implementation uses **manual gradient descent** whereas sklearn implementation uses a **different solver (LBFGS)**, **which converges differently and can reach slightly different optima on the same dataset.**
- There are also various other parameters (tolerance, max_iters) that sklearn sets that we were unable to account for with our basic implementation.
- Either way we were able to demonstrate **highly comparable results to a slight degree of divergence.**

Summary - what was interesting/challenging?

1. **Manually tuning** hyperparameters was difficult: (e.g., finding the right combination of things like batch size and lambda)
2. **Preprocessing** the data into a suitable representation was also fairly involved (e.g., scaling them, encoding labels)
3. **Vectorizing** the pseudocode in Numpy and making sure shapes were compatible

References

Khodabakhsh, Hojjat. *MNIST Dataset*. Kaggle. Accessed December 10, 2025.
<https://www.kaggle.com/datasets/hojjatk/mnist-dataset>.

Pedregosa, F. et al., 2011. Scikit-learn: Machine learning in Python. *Journal of machine learning research*, 12(Oct), pp.2825–2830.

Shalev-Shwartz, Shai, and Shai Ben-David. *Understanding Machine Learning: From Theory to Algorithms*. Cambridge university press, 2014.

Zsom, Andras. “Lecture 5: Logistic Regression.”, Lecture, Brown University, September 18, 2025.

Zsom, Andras. “Lecture 6: SGD, Data Prep, and other Practicalities.”, Lecture, Brown University, September 23, 2025.